Vinod Sagar - Detailed CV Summary

AI Engineer | Generative AI | Agentic AI | RAG Pipelines | LLM Systems

Bengaluru, India | Mobile: +91 8660687039 | vinodsagar.aieng@gmail.com | linkedin.com/in/vinod-sagar-b504b6111 | github.com/vinodsagar260981

Profile Summary

Vinod Sagar is an AI Engineer based in Bengaluru, India with a strong focus on Generative AI, Agentic AI, LLM systems, and RAG pipelines. His profile highlights practical experience in building production-grade AI systems using LangGraph, LangChain, MCP, Microsoft Foundry, FastAPI, Django, vector databases, and AWS-based deployments.

He combines a traditional software engineering background in Python, Django, APIs, and data systems with modern AI engineering skills such as LLM integration, prompt engineering, multi-agent workflows, retrieval systems, model evaluation, and secure AI application development.

Experience Summary

Total experience across current and previous software/AI companies is about 7 years, calculated from HCL Technologies (August 2018 to July 2023) and Tap Health (August 2024 to present, as of August 2026).

Relevant experience in AI engineering is around 3 years, covering Generative AI, Agentic AI, RAG architecture, LLM application development, LangChain, LangGraph, FastAPI, MCP, Microsoft Foundry, vector databases, and production-oriented AI workflows.

Core AI Skills

The CV highlights strong expertise in Agentic AI frameworks and protocols, including LangGraph workflows, LangChain agents, retrievers, vector stores, LCEL, callback systems, and MCP-based client-server architecture.

His AI engineering skill set includes LLM gateways, AI guardrails, PII redaction, Human-in-the-Loop validation, LangSmith-based evaluation, hallucination detection, golden dataset regression testing, model routing, fallback model handling, and cost and latency optimization.

He has worked with major LLM ecosystems such as OpenAI GPT models, Meta Llama, Mistral, Anthropic Claude, Google Gemini, Groq, Kimi, and Hugging Face models. His model optimization knowledge includes LoRA, QLoRA, PEFT, embedding optimization, Chain-of-Thought prompting, ReAct prompting, few-shot prompting, and retrieval-based context augmentation.

Technical Stack

The CV lists hands-on experience with vector and search technologies such as Pinecone, FAISS, Chroma, semantic search, hybrid search, metadata filtering, similarity search, contextual retrieval, and document chunking strategies for RAG systems.

Backend and API experience includes FastAPI, Django, Django REST Framework, Pydantic, REST APIs, JWT authentication, OAuth2, asynchronous services, and AI model inference endpoints.

Cloud and DevOps exposure includes AWS Bedrock, AWS Lambda, EC2, S3, CloudWatch, Docker, Docker Compose, and containerized AI deployments.

Current Role - Tap Health

Vinod is currently working as an AI Engineer at Tap Health in Bengaluru from August 2024 to present. Tap Health is described as a startup with around 15 team members, giving him broad exposure to end-to-end AI product development in a fast-moving environment.

His role includes designing and deploying production-grade RAG systems, building multi-agent workflows with LangGraph, integrating MCP for structured context sharing, developing Microsoft Foundry AI agents, implementing Human-in-the-Loop checkpoints, securing APIs with FastAPI and JWT/OAuth2, fine-tuning domain-specific LLMs, and deploying workloads on AWS.

Project Experience

The Agentic AI Workflow System project uses LangGraph, MCP, FastAPI, and Docker. It includes multi-agent architecture, tool nodes, subgraphs, persistence, memory, HITL checkpoints, secure APIs, and Docker-based deployment.

The Generative AI Agent project uses Microsoft Foundry, FastAPI, and RAG. It focuses on model orchestration, contextual responses, vector database integration, modular tool calling, and multi-step reasoning.

The Conversational RAG Chatbot project uses LangChain, Gemini, Pinecone, and FastAPI. It includes multi-turn conversation memory, semantic and keyword retrieval, hybrid search, and API deployment.

The Secure AI API Platform project includes FastAPI, Django, LangGraph, MCP, LLM Gateway, Descope, Docker, and AWS. It covers RBAC, secure authorization, model routing, fallback models, rate limiting, guardrails, PII handling, evaluation workflows, hallucination detection, regression testing, and AWS monitoring.

Previous Experience

Before Tap Health, Vinod worked at HCL Technologies from August 2018 to July 2023 as a Software Engineer. His work included Python development at Google India, data engineering at Smith & Nephew, Django development for Smart DCC in Manchester and Bengaluru, and Python automation for Caterpillar.

From August 2013 to July 2018, he worked as a CAD Customization Engineer for organizations including Daimler, Volvo, Stelia Aerospace, and Mercedes-Benz R&D; India, mainly focusing on CAD automation using VB.NET.

Career Break

The CV explains a 13-month career break from July 2023 to August 2024 due to personal reasons. During this period, Vinod remained professionally active by upgrading his skills in Generative AI, Agentic AI, RAG architecture, LLM application development, prompt engineering, vector databases, LangChain, LangGraph, FastAPI, MCP, Microsoft Foundry, Pinecone, AWS Bedrock, and production-oriented AI workflows.

The break is presented as a constructive phase that helped him transition from a Python/Django software engineering background into modern AI engineering.

Compensation, Availability, and Education

The CV states that his current compensation is 10 LPA, expected compensation is 30 LPA, and he can join within 15 days.

His education is Bachelor of Engineering in Mechanical Engineering from Bangalore Institute of Technology, Bengaluru, completed in 2008. His listed interests include wildlife photography, solo bike riding, trekking, cooking, and farming. Languages listed are English, Kannada, Hindi, Telugu, and Tamil.